

COMPUTER SYSTEMS AND ORGANIZATION

Sockets

Daniel G. Graham Ph.D




UNIVERSITY
of VIRGINIA

ENGINEERING




1. Client-server model
2. HTTP protocol basic
3. TCP Client
4. Client-server example demo
5. System Calls

CLIENT SERVER MODEL



Two types of sockets

CLIENT SERVER MODEL



Two types of sockets

DNS: FINDING THE IP FOR A DOMAIN

```
dgg6b@Daniels-Mac-mini ~ % dig bing.com
```

```
; <<>> DiG 9.10.6 <<>> bing.com
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 54193
;; flags: qr rd ra; QUERY: 1, ANSWER: 2, AUTHORITY: 0, ADDITIONAL: 1

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; udp: 512
;; QUESTION SECTION:
;bing.com.                IN      A

;; ANSWER SECTION:
bing.com.                 1451    IN      A      13.107.21.200
bing.com.                 1451    IN      A      204.79.197.200

;; Query time: 28 msec
;; SERVER: 8.8.8.8#53(8.8.8.8)
```

Use dig to look up
IP address for a
particular website

HTTP BASICS

REQUEST

```
GET /index HTTP/1.1\r\n
Host: www.bing.com\r\n
\r\n
```

Response

```
HTTP/1.1 200 OK
--- Headers ---
--- Content ----
```

line wrap

```
1 <!doctype html><html lang="en" dir="ltr"><head><
2   style="position:relative;vertical-align:top;mar
3   var _d,sb_de;typeof _d=="undefined"&&(_d=document
4   //]]></script><a id="id_mobile" class="id_button
5   var img_p = document.getElementById('id_p'); img
6   //]]></script><script type="text/javascript" non
7   var preloadBg = document.getElementById('preload
8   //]]></script><script type="text/javascript" cro
9   0;function getBrowserWidth_Desk(){var t=_d.docum
10  //]]></script><script type="text/javascript" cro
11  sa_config={"f":"sb_form","i":"sb_form_q","c":"sw
12  //]]></script><div id="aRmsDefer"><script type="
13  var mcp_banner=function(n){function u(n){var t=s
14  //]]></script><script type="text/rms" nonce="w5i
15  0;
16  //]]></script><script type="text/rms" nonce="w5i
17  if (typeof(PrefetchJsResource) !== "undefined")
18  //]]></script><script type="text/rms" nonce="w5i
19  var sj_appHTML=function(n,t){var f,e,o,r,i,s,h;i
20  //]]></script><script type="text/rms" nonce="w5i
21  Feedback.Bootstrap.InitializeFeedback({page:true
22  //]]></script><script type="text/rms" nonce="w5i
23  _G!==undefined&&_G.EF!==undefined&&_G.EF.bmasync
24  //]]></script><script type="text/rms" nonce="w5i
```



Search the web

Create a table that analyzes
the arts compared to



Ask Bing Chat

NOW LET'S WRITE A PROGRAM

Let's write a c program that will send an HTTP request to the Bing servers and get the index page.

PART 1

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>

#define PORT 80
#define BUFFER_SIZE 4096
#define SERVER_IP "13.107.21.200"

int main() {
    int sock;
    struct sockaddr_in server;
    char message[BUFFER_SIZE], response[BUFFER_SIZE];

    // Create socket
    sock = socket(AF_INET, SOCK_STREAM, 0);

    // Prepare the sockaddr_in structure
    server.sin_addr.s_addr = inet_addr("SERVER_IP");
    server.sin_family = AF_INET;
    server.sin_port = htons(PORT);

    // Connect to the server
    connect(sock, (struct sockaddr *)&server, sizeof(server));
```

Client vs Server:

Notice that we use
connect instead of
accept.

PART 2

```
// Create GET request
snprintf(message, sizeof(message), "GET / HTTP/1.1\r\nHost: www.bing.com\r\n\r\n");


// Send the message
write(sock, message, strlen(message));

// Receive the server's response
read(sock, response, BUFFER_SIZE);

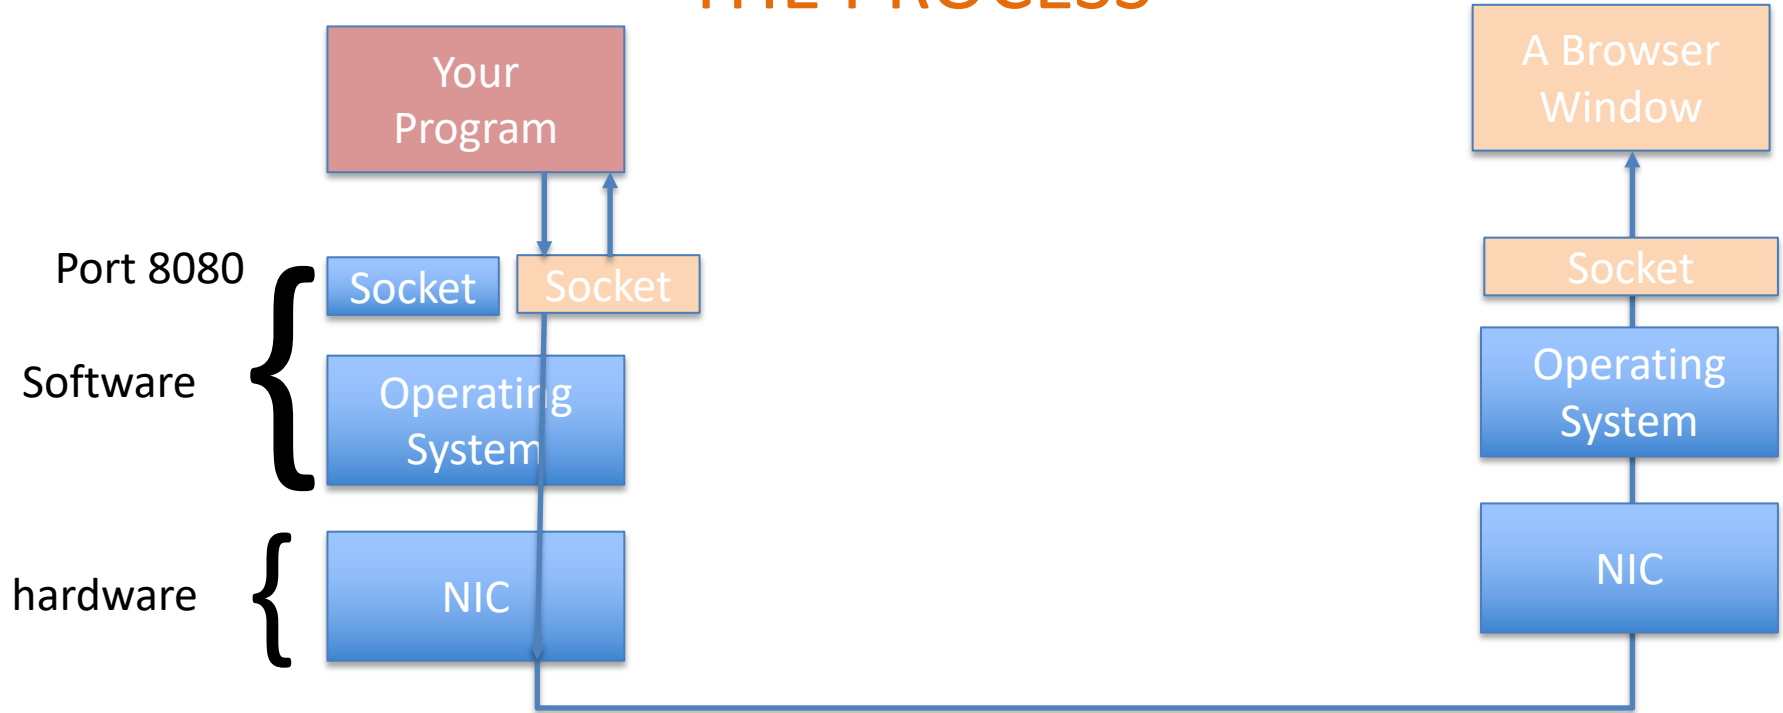
printf("Server Response:\n%s", response);

// Close the socket
close(sock);

return 0;
}
```



THE PROCESS



OUR SERVER

We have implemented both the client and server.

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <netinet/in.h>

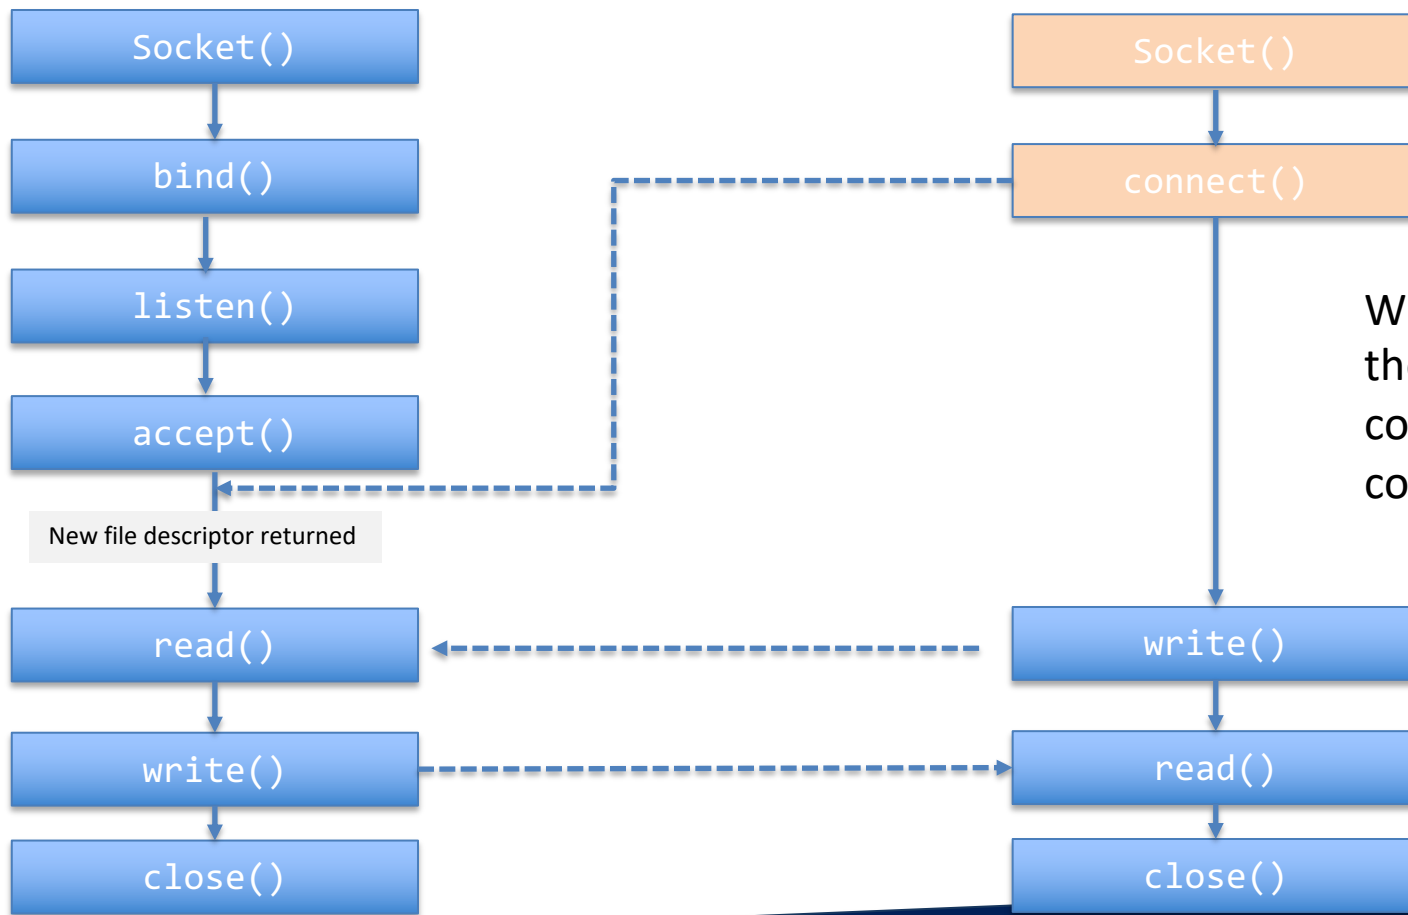
#define PORT 8080

int main() {
    int server_fd, client_fd;
    struct sockaddr_in address;

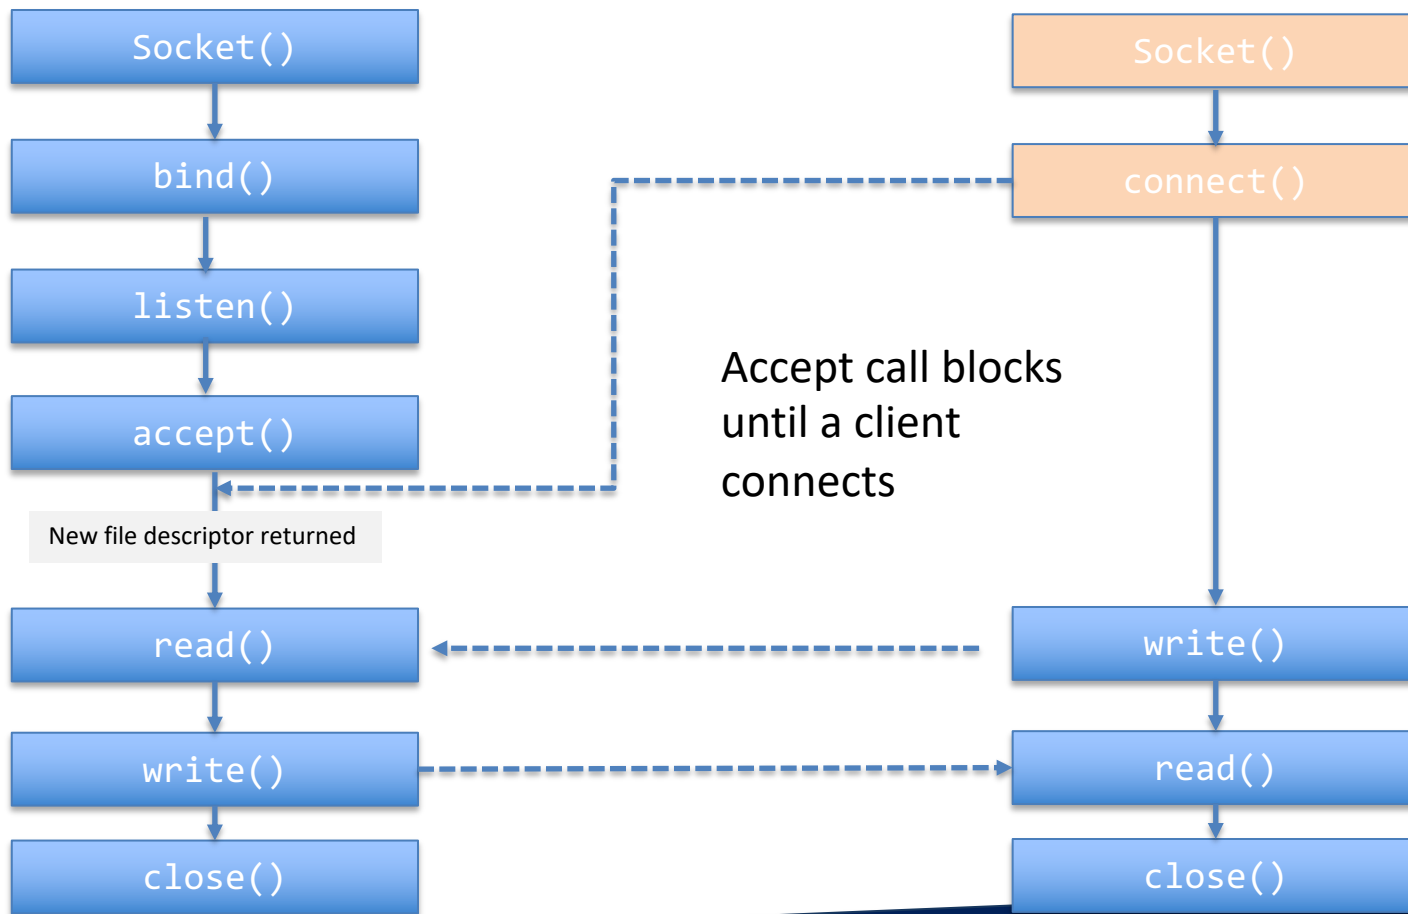
    server_fd = socket(AF_INET, SOCK_STREAM, 0);
    address.sin_family = AF_INET;
    address.sin_addr.s_addr = INADDR_ANY;
    address.sin_port = htons(PORT);

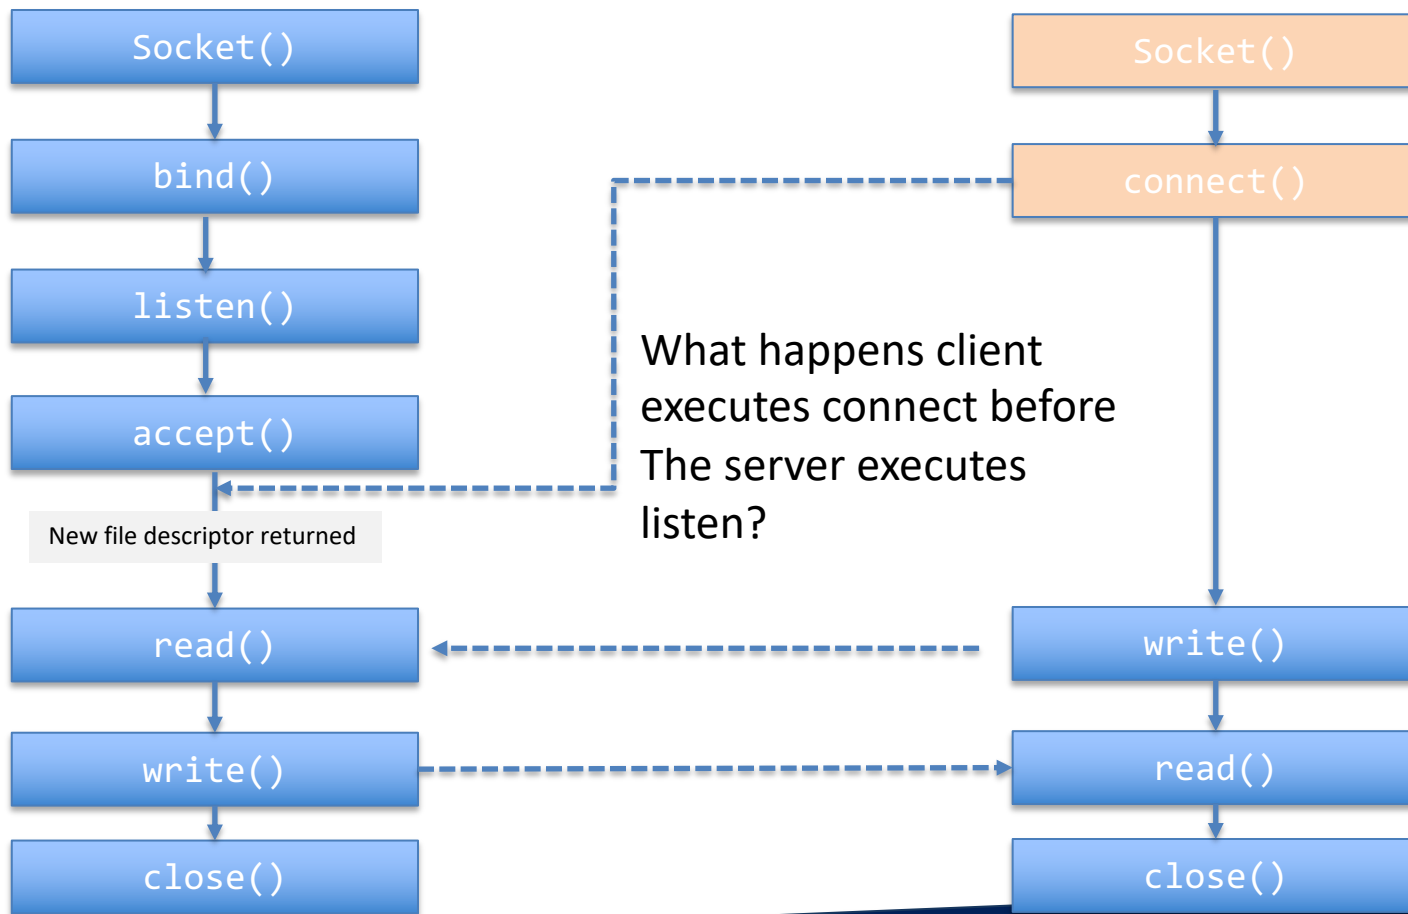
    bind(server_fd, (struct sockaddr *)&address, sizeof(address));
    listen(server_fd, 10);
    int addrlen = sizeof(address);

    while (1) {
        new_socket = accept(server_fd, (struct sockaddr *)&address, (socklen_t*)&addrlen);
        write(new_socket, "HTTP/1.1 200 OK\n", 16);
        write(new_socket, "Content-Type: text/html\n\n", 25);
        write(new_socket, "<html><body><h1>Hello, World!</h1></body></html>", 44);
        close(new_socket);
    }
    close(server_fd);
    return 0;
}
```



What happens if
the client get to
connect a section of
code before get t





WE HAVE BEEN USING FUNCTIONS LIKE WRITE HOW DOES THAT GET IMPLEMENTED IN ASSEMBLY?

```
write(new_socket, "HTTP/1.1 200 OK\n", 16);
```

```
#include <unistd.h>
#include <fcntl.h>

int main() {
    int fd;
    char *text = "CS01";

    // Open a file for writing (create it if it doesn't exist)
    fd = open("output.txt", O_WRONLY | O_CREAT, 0644);

    // Write the string to the file
    write(fd, text, 4); // 4 is the number of bytes to write

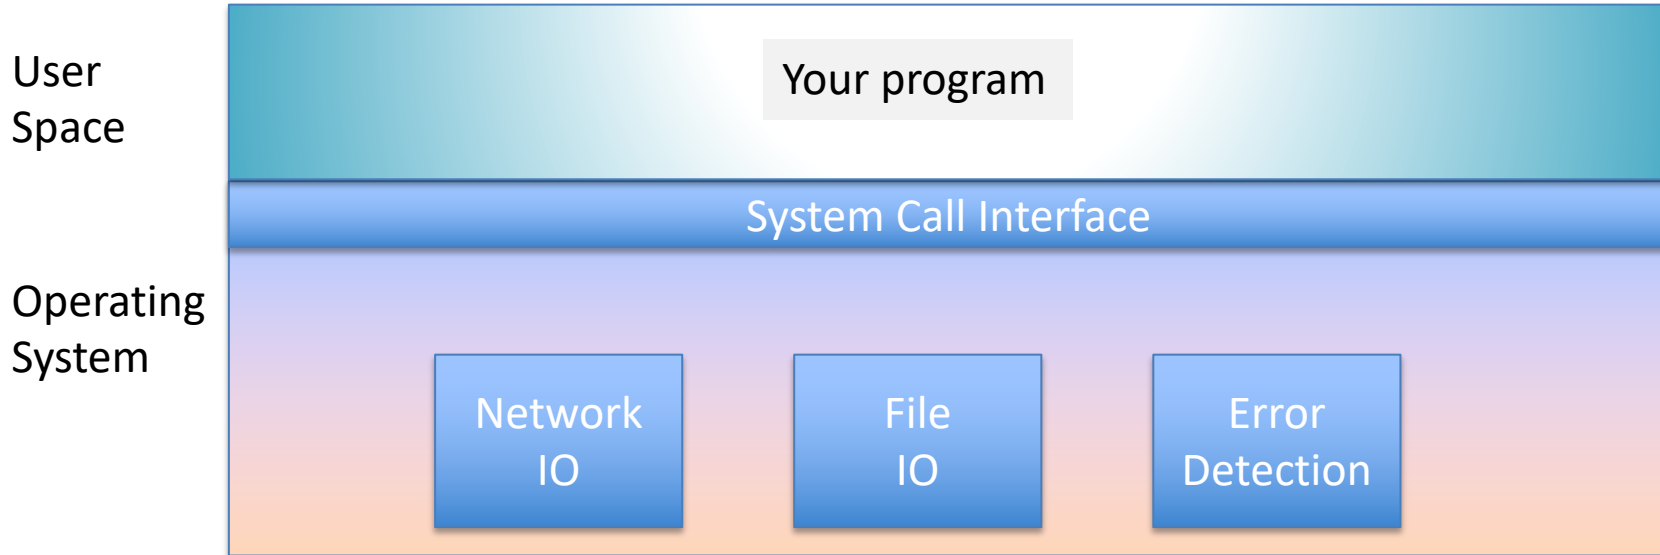
    // Close the file
    close(fd);

    return 0;
}
```

WHAT DOES THIS LOOK LIKE IN ASSEMBLY?

Let's look at this one.
Sadly it not simple a
call instruction to
function located in
fcntl

READING AND WRITING FILES AND THE NETWORK IS A PRIVILEGED OPERATION



USER SPACE VS KERNEL SPACE LINUX

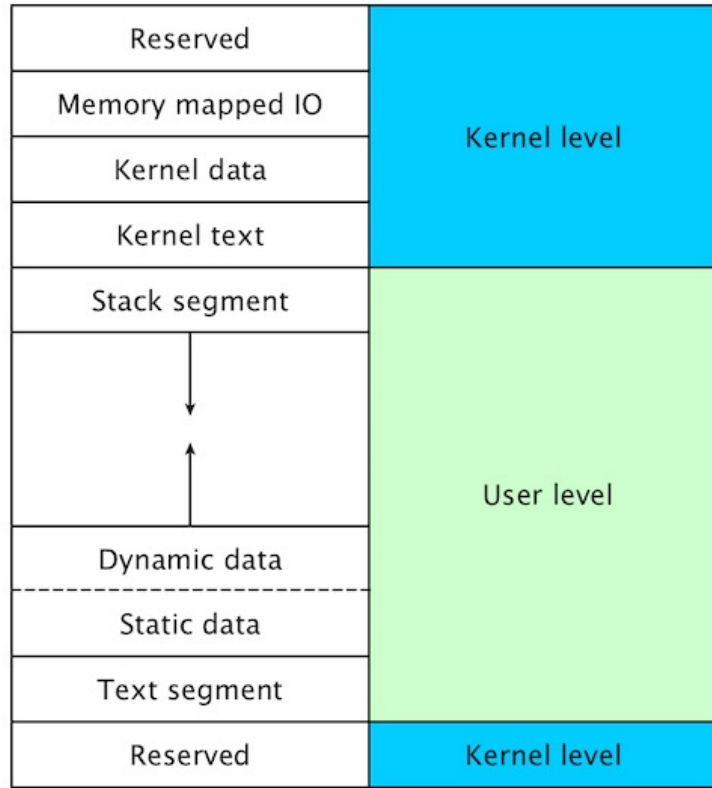
0xffffffff

0xfffff0010

0xfffff0000

0x90000000

0x80000000



0x10000000

0x04000000

0x00000000

Kernel layout for MIPS chips

<https://www.it.uu.se/education/course/homepage/os/vt18/module-0/mips-and-mars/mips-memory-layout/>

The layout of the arm chips can be found here.

<https://www.kernel.org/doc/html/v5.7/arm/memory.html>

SYSTEM CALL CALLING CONVENTION

1. Register Usage for Arguments:

1. `%rax`: System call number. Each system call has a unique number that you place in this register to tell the kernel which system call you're making.
2. `%rdi`, `%rsi`, `%rdx`, `%r10`, `%r8`, `%r9`: Used for passing up to six arguments to system calls. `%rdi` is for the first argument, `%rsi` for the second, and so on. If a system call needs more than six arguments, a pointer to a block containing the arguments is passed as one of these registers.

2. Making the System Call:

1. The `syscall` instruction is used to switch to kernel mode and invoke the system call. The kernel examines the value in `%rax` and understands which system call is being requested.

3. Return Value:

1. After the system call, the return value is placed in `%rax`. This value typically indicates success or an error code.

THING ABOUT HOW YOU IMPLEMENT THE WRITE SYSTEM CALL TO STDOUT

```
write(1, message, message_length)
```

1. Register Usage for Arguments:

1. `%rax`: System call number. Each system call has a unique number that you place in this register to tell the kernel which system call you're making.
2. `%rdi`, `%rsi`, `%rdx`, `%r10`, `%r8`, `%r9`: Used for passing up to six arguments to system calls. `%rdi` is for the first argument, `%rsi` for the second, and so on. If a system call needs more than six arguments, a pointer to a block containing the arguments is passed as one of these registers.

2. Making the System Call:

1. The `syscall` instruction is used to switch to kernel mode and invoke the system call. The kernel examines the value in `%rax` and understands which system call is being requested.

3. Return Value:

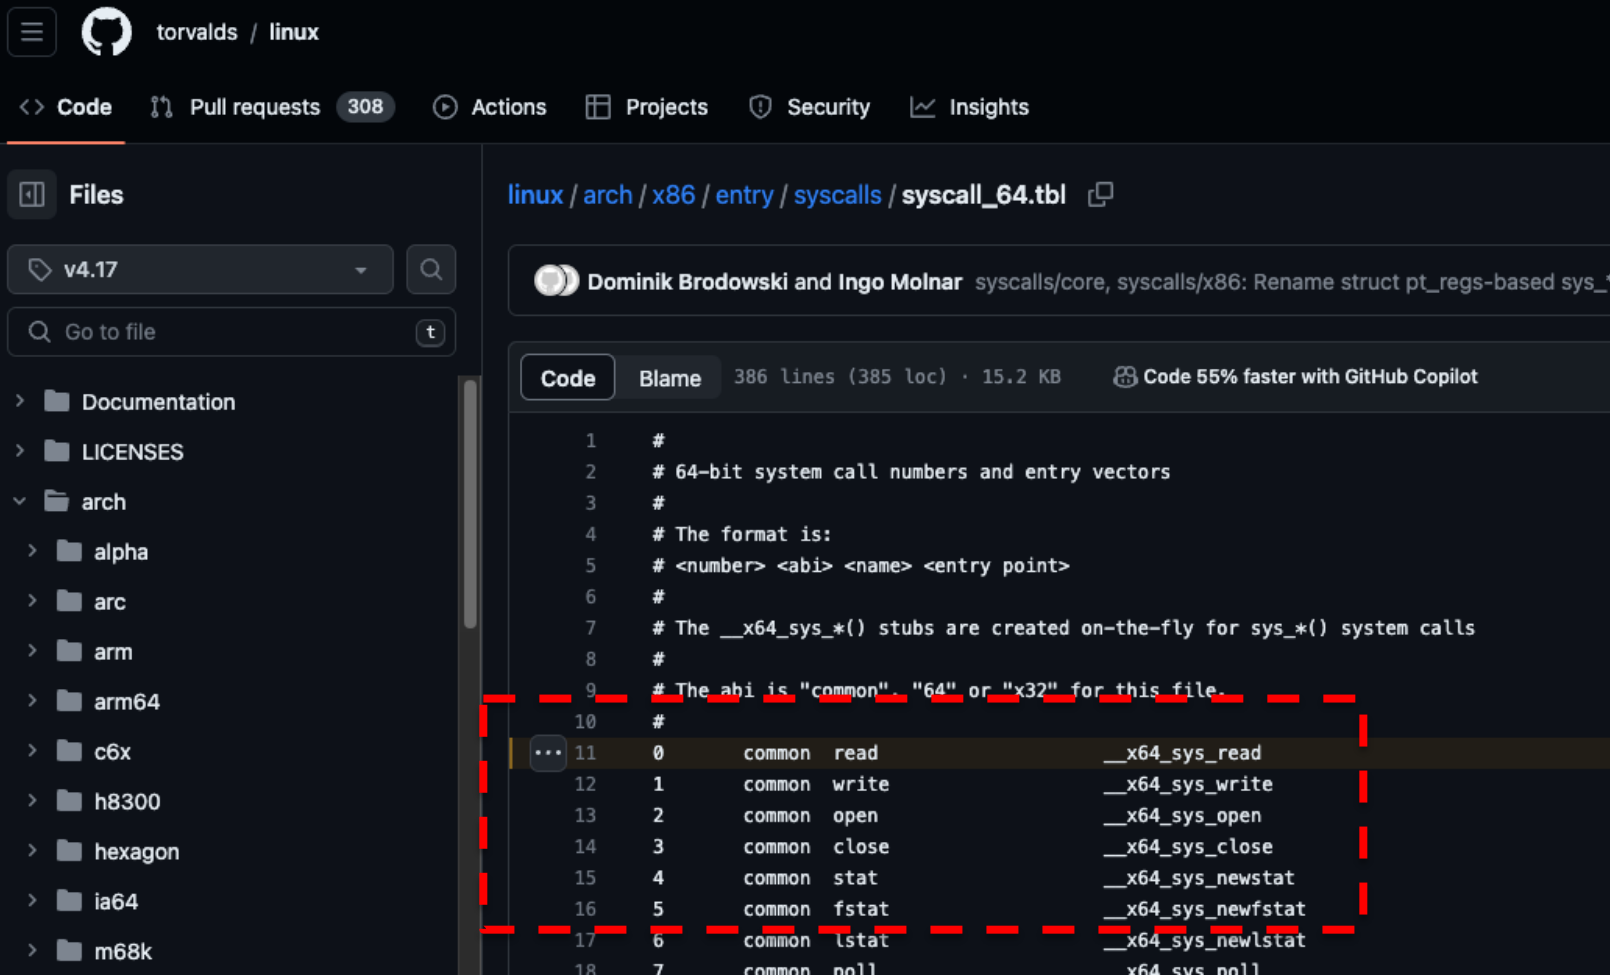
1. After the system call, the return value is placed in `%rax`. This value typically indicates success or an error code.

SYSTEM CALL CALLING CONVENTION

```
.global _start
.text
_start:
    # write(1, message, 18)
    mov     $1, %rax           ; syscall number for write (1)
    mov     $1, %rdi           ; file descriptor 1 (stdout)
    lea     message(%rip), %rsi ; load the address of the message
    mov     $18, %rdx          ; message length (18 bytes)
    syscall                    ; perform the system call

.section .rodata               ; Read-only data section
message:                       ; Label for the message
    .ascii "Computer Systems 1" ;
```

WHERE CAN I FIND THE SYSTEM CALL NUMBERS



The screenshot shows the GitHub interface for the Linux kernel repository. The file path is `linux/arch/x86/entry/syscalls/syscall_64.tbl`. The file is 15.2 KB and contains 386 lines of code. The code is displayed in a table format, showing system call numbers, their ABI, names, and entry points. A red dashed box highlights the first 10 lines of the table, which are the 64-bit system call numbers and entry vectors.

Files

v4.17

Go to file

Documentation

LICENSES

arch

alpha

arc

arm

arm64

c6x

h8300

hexagon

ia64

m68k

linux / arch / x86 / entry / syscalls / syscall_64.tbl

Dominik Brodowski and Ingo Molnar syscalls/core, syscalls/x86: Rename struct pt_regs-based sys_*

Code Blame 386 lines (385 loc) · 15.2 KB Code 55% faster with GitHub Copilot

```
1 #
2 # 64-bit system call numbers and entry vectors
3 #
4 # The format is:
5 # <number> <abi> <name> <entry point>
6 #
7 # The __x64_sys_*() stubs are created on-the-fly for sys_*() system calls
8 #
9 # The abi is "common", "64" or "x32" for this file.
10 #
11 0 common read __x64_sys_read
12 1 common write __x64_sys_write
13 2 common open __x64_sys_open
14 3 common close __x64_sys_close
15 4 common stat __x64_sys_newstat
16 5 common fstat __x64_sys_newfstat
17 6 common lstat __x64_sys_newlstat
18 7 common poll x64_sys_poll
```

Linux github
repo.

https://github.com/torvalds/linux/blob/v4.17/arch/x86/entry/syscalls/syscall_64.tbl

SYSTEM CALLS

```
int main() {  
    int fd;  
    char *text = "CS01";  
  
    // Open a file for writing (create it if it doesn't exist)  
    fd = open("output.txt", O_WRONLY | O_CREAT, 0644)
```

User space

Returns file descriptor

System Call Interface

Kernel space

Call #	Function pointer
0	read
1	write
2	open
3	close

open()
implementation of open
file descriptor setup etc

return

WHAT DOES THE FOLLOWING ASSEMBLY DO?

```
.global _start
.text
_start:
    # What does this snippet of assemble do?
    mov     $3, %rax          ;
    mov     $1, %rdi          ;
    syscall                   ;
```

Call #	Function pointer
0	read
1	write
2	open
3	close

- A. Write Perror
- B. Write stdout
- C. Open stdout
- D. Open Perror
- E. Read from Perror
- F. Close Perror
- G. Read stdout
- H. Close stdout
- I. Read stdin
- J. Close std in

